

Project 1 - California Water Usage

Welcome to your first project! We will be exploring possible connections between water usage, geography, and income in California. The water data for this project was procured from the [California State Water Resources Control Board](#) and curated by the [Pacific Institute](#). The map data includes [US topography](#), [California counties](#), and [ZIP codes](#).

The dataset on income comes from the IRS ([documented here](#)). We have identified some interesting columns in the dataset, but a full description of all the columns (and a definition of the population in the dataset and some interesting anonymization procedures they used) is available [here](#).

Due Date and Comments

Your project will be due **Friday, March 3 at 11PM**. If you like, you may work with one partner. Outside of your partner, your work should be your own work. **This project has some hard pieces!** It is okay if you do not solve every task, but you should try hard on every one and document your efforts. I and the TAs (Zoe and Lesley) are very happy to help with hints and nudges as well as looking for bugs. If you start early, you should have plenty of time to ask questions.

```
In [1]: # Run this cell, but please don't change it.

import numpy as np
import math
from datascience import *

# These lines set up the plotting functionality and formatting.
import matplotlib
matplotlib.use('Agg')
%matplotlib inline
import matplotlib.pyplot as plots
plots.style.use('fivethirtyeight')
```

First, load the data. Loading may take some time. **You will need to upload the two files from Moodle to the same location on your Jupyter hub that you have Project 1 located.** First download them from Moodle and then use the Upload button on the Jupyter hub to upload them.

```
In [2]: # Run this cell, but please don't change it.

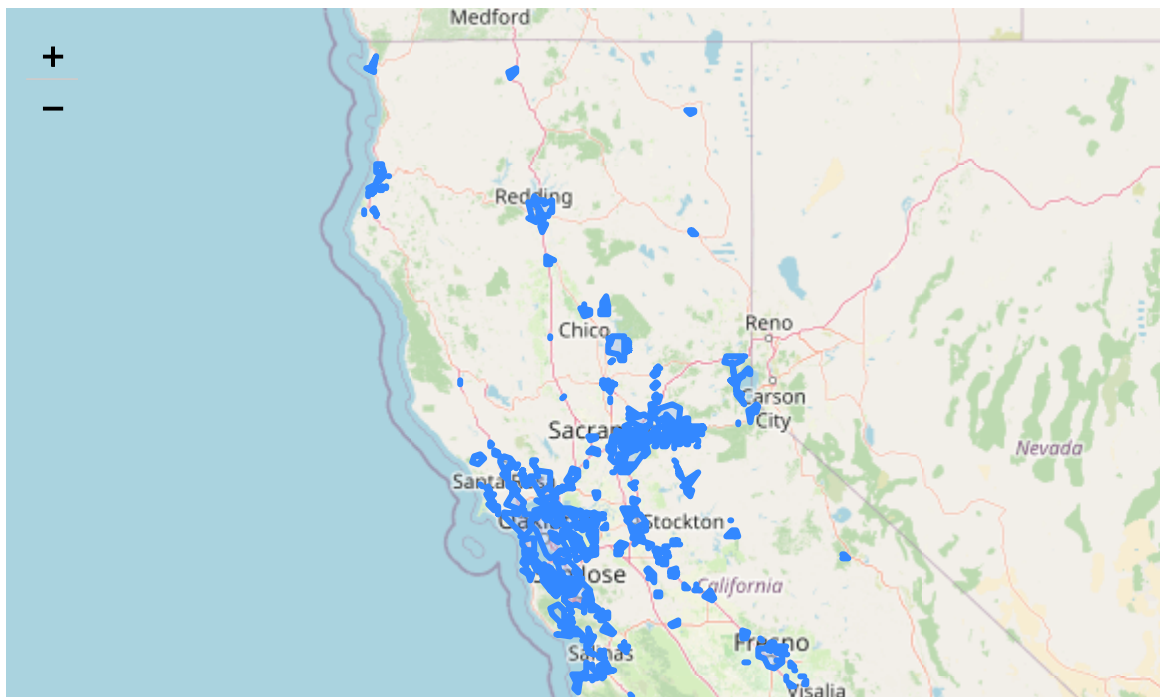
districts = Map.read_geojson('water_districts.geojson')
zips = Map.read_geojson('ca_zips.geojson.gz')
usage_raw = Table.read_table('~/.DS_113_S23/Projects/Project_1/water_usage.csv')
income_raw = Table.read_table('~/.DS_113_S23/Projects/Project_1/ca_income_by_zip.csv')
wd_vs_zip = Table.read_table('~/.DS_113_S23/Projects/Project_1/wd_vs_zip.csv')
```

Part 1: Maps

The `districts` and `zips` data sets are `Map` objects. Documentation on mapping in the `datascience` package can be found at data8.org/datascience/maps.html. To view a map of California's water districts, run the cell below. Click on a district to see its description.

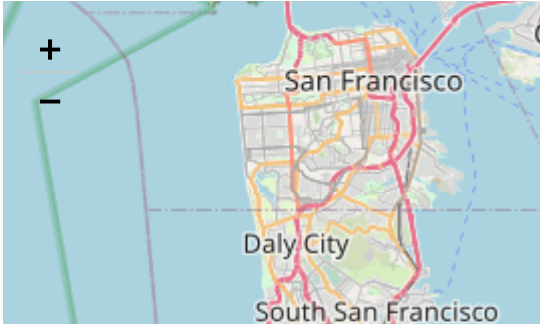
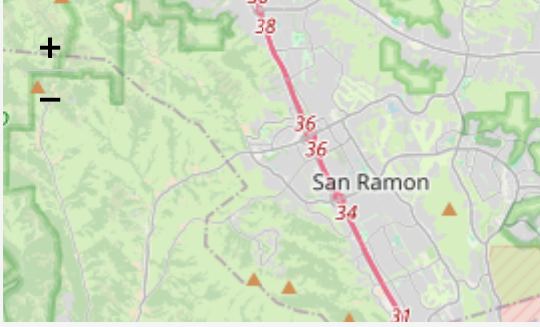

```
In [3]: districts.format(width=800, height=600)
```

Out[3]:



A `Map` is a collection of regions and other features such as points and markers, each of which has a **string** `id` and various properties. You can view the features of the `districts` map as a table using `Table.from_records`.

```
In [4]: district_table = Table.from_records(districts.features)
district_table.show(3)
```

| PWSID | feature | id | popupContent |
|---------|--|----|---|
| 0110001 |  | 0 | Alameda County Water District |
| 0110003 |  | 1 | California Water Service Company Livermore |
| 0110005 |  | 2 | East Bay Municipal Utilities District |

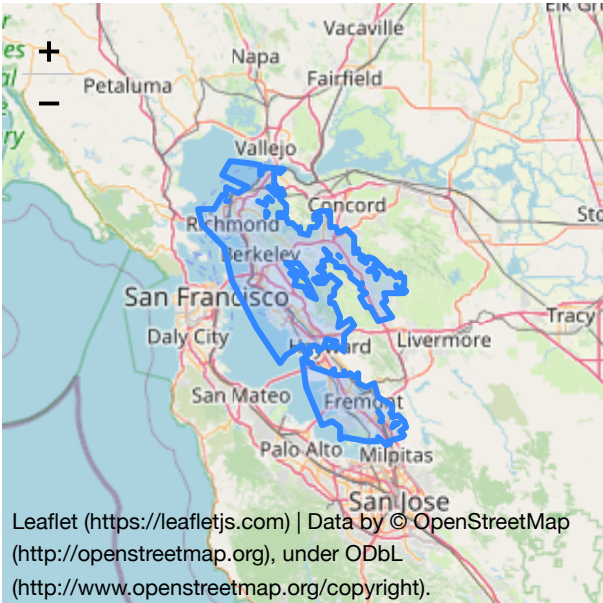
... (407 rows omitted)

To display a `Map` containing only two features from the `district_table`, call `Map` on an array containing those two features from the `feature` column.

Question 1.1. Draw a map of the Alameda County Water District (row 0) and the East Bay Municipal Utilities District (row 2).

```
In [5]: # Fill in the next line to make an array containing the two elements from d1
# that you want out of the feature column.
alameda_and_east_bay = make_array(district_table["feature"][0],district_table["feature"][2])
Map(alameda_and_east_bay, height=300, width=300)
```

Out [5] :



> Trust Notebook

Hint: If scrolling becomes slow on your computer, you can clear maps for the cells above by running `Cell > All Output > Clear` from the `Cell` menu.

Part 2: California Income

Let's look at the `income_raw` table, which comes from the IRS. We're going to link this information about incomes to our information about water. First, we need to investigate the income data and get it into a more usable form.

In [6]: `income_raw`

Out [6]:

| | ZIP | N1 | MARS1 | MARS2 | MARS4 | PREP | N2 | NUMDEP | A00100 | N02650 | A02650 |
|--|-------|-------|-------|-------|-------|-------|-------|--------|--------|--------|--------|
| | 90001 | 13100 | 6900 | 1890 | 4270 | 10740 | 29670 | 15200 | 181693 | 13100 | 184344 |
| | 90001 | 5900 | 1700 | 1970 | 2210 | 4960 | 17550 | 9690 | 203628 | 5900 | 204512 |
| | 90001 | 1480 | 330 | 760 | 390 | 1240 | 4710 | 2470 | 89065 | 1480 | 89344 |
| | 90001 | 330 | 50 | 210 | 70 | 290 | 1100 | 560 | 28395 | 330 | 28555 |
| | 90001 | 160 | 30 | 100 | 40 | 130 | 510 | 250 | 24676 | 160 | 25017 |
| | 90001 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 90002 | 12150 | 6330 | 1460 | 4330 | 9580 | 27240 | 14070 | 167261 | 12150 | 170095 |
| | 90002 | 5030 | 1510 | 1490 | 1980 | 4120 | 14410 | 7890 | 173280 | 5030 | 174335 |
| | 90002 | 1320 | 300 | 600 | 400 | 1060 | 4090 | 2180 | 78559 | 1320 | 78871 |
| | 90002 | 340 | 90 | 190 | 90 | 270 | 1060 | 530 | 28502 | 340 | 28558 |

... (8888 rows omitted)

Some observations:

1. The table contains several numerical columns and a column for the ZIP code.

- For each ZIP code, there are 6 rows. Each row for a ZIP code has data from tax returns in one *income bracket*. (A tax return is the tax filing from one person or household. An income bracket is a group of people whose annual income is in some range, like 25,000 USD to 34,999 USD.)
- According to the IRS documentation, all the numerical columns are *totals* -- either total numbers of returns that fall into various categories, or total amounts of money (in thousands of dollars) from returns in those categories. For example, the column `'N02650'` is the number of returns that included a total income amount, and `'A02650'` is the total amount of total income (in thousands of dollars) from those returns.

For the analysis we're about to do, we won't need to use the information about tax brackets. We will need to know the total income, total number of returns, and other totals from each ZIP code.

Question 2.1. Assign the name `income_by_zipcode` to a table with just one row per ZIP code. When you group according to ZIP code, the remaining columns should be summed. In other words, for any other column such as `'N02650'`, the value of `'N02650'` in a row corresponding to ZIP code 90210 (for example) should be the sum of the values of `'N02650'` in the 6 rows of `income_raw` corresponding to ZIP code 90210.

In [7]:

```
income_by_zipcode = income_raw.groupby("ZIP", np.sum)
income_by_zipcode
```

Out[7]:

| ZIP | N1 sum | MARS1 sum | MARS2 sum | MARS4 sum | PREP sum | N2 sum | NUMDEP sum | A00100 sum | N02650 sum |
|-------|-----------|--------------|--------------|--------------|-------------|-----------|---------------|---------------|---------------|
| 90001 | 20970 | 9010 | 4930 | 6980 | 17360 | 53540 | 28170 | 527457 | 20970 |
| 90002 | 18960 | 8230 | 3830 | 6800 | 15120 | 47200 | 24850 | 462823 | 18960 |
| 90003 | 26180 | 11310 | 5130 | 9640 | 20570 | 64470 | 33760 | 612733 | 26180 |
| 90004 | 27360 | 15330 | 7000 | 4670 | 20260 | 51180 | 17800 | 1.61777e+06 | 27360 |
| 90005 | 15430 | 8550 | 3870 | 2830 | 11210 | 29910 | 11130 | 707020 | 15430 |
| 90006 | 22630 | 11470 | 5400 | 5630 | 17840 | 47590 | 20210 | 563530 | 22630 |
| 90007 | 11710 | 6350 | 2270 | 3020 | 8310 | 23380 | 9950 | 311779 | 11710 |
| 90008 | 14710 | 8060 | 2310 | 4110 | 9990 | 27000 | 10310 | 662036 | 14710 |
| 90010 | 2210 | 1270 | 690 | 210 | 1760 | 3790 | 960 | 314333 | 2210 |
| 90011 | 36670 | 15540 | 8600 | 12390 | 30240 | 95640 | 51260 | 857731 | 36670 |

... (1473 rows omitted)

Your `income_by_zipcode` table probably has column names like `N1 sum`, which looks a little weird.

Question 2.2. Relabel the columns in `income_by_zipcode` to match the labels in `income_raw`

Hint: Inspect `income_raw.labels` and `income_by_zipcode.labels` to find the differences you need to change.

Hint 2: Since there are many columns, it will be easier to relabel each of them by using a `for` statement. See Chapter 8 of the textbook for details.

Hint 3: You can use the `replace` method of a string to remove excess content. See [lab02](#) for examples.

Hint 4: To create a new table from an existing table with one label replaced, use `relabelled`. To **change** a label in an existing table permanently, use `relabel`. Both methods take two arguments: the old label and the new label. You can solve this problem with either one, but `relabel` is simpler.

```
In [8]: for x in income_by_zipcode.labels:
        income_by_zipcode.relabel(x,x.replace(" sum",""))
        income_by_zipcode
```

```
Out[8]:
```

| | ZIP | N1 | MARS1 | MARS2 | MARS4 | PREP | N2 | NUMDEP | A00100 | N02650 |
|--|-------|-------|-------|-------|-------|-------|-------|--------|-------------|--------|
| | 90001 | 20970 | 9010 | 4930 | 6980 | 17360 | 53540 | 28170 | 527457 | 20970 |
| | 90002 | 18960 | 8230 | 3830 | 6800 | 15120 | 47200 | 24850 | 462823 | 18960 |
| | 90003 | 26180 | 11310 | 5130 | 9640 | 20570 | 64470 | 33760 | 612733 | 26180 |
| | 90004 | 27360 | 15330 | 7000 | 4670 | 20260 | 51180 | 17800 | 1.61777e+06 | 27360 |
| | 90005 | 15430 | 8550 | 3870 | 2830 | 11210 | 29910 | 11130 | 707020 | 15430 |
| | 90006 | 22630 | 11470 | 5400 | 5630 | 17840 | 47590 | 20210 | 563530 | 22630 |
| | 90007 | 11710 | 6350 | 2270 | 3020 | 8310 | 23380 | 9950 | 311779 | 11710 |
| | 90008 | 14710 | 8060 | 2310 | 4110 | 9990 | 27000 | 10310 | 662036 | 14710 |
| | 90010 | 2210 | 1270 | 690 | 210 | 1760 | 3790 | 960 | 314333 | 2210 |
| | 90011 | 36670 | 15540 | 8600 | 12390 | 30240 | 95640 | 51260 | 857731 | 36670 |

... (1473 rows omitted)

Question 2.3. Create a table called `income` with one row per ZIP code and the following columns.

1. A `ZIP` column with the same contents as `'ZIP'` from `income_by_zipcode`.
2. A `num returns` column containing the total number of tax returns that include a total income amount (column `'N02650'` from `income_by_zipcode`).
3. A `total income ($)` column containing the total income in all tax returns in thousands of dollars (column `'A02650'` from `income_by_zipcode`).
4. A `num farmers` column containing the number of farmer returns (column `'SCHF'` from `income_by_zipcode`).

```
In [9]: income = Table().with_columns(
        "ZIP", income_by_zipcode["ZIP"],
        "num returns", income_by_zipcode["N02650"],
        "total income ($)", income_by_zipcode["A02650"],
        "num farmers", income_by_zipcode["SCHF"]
    )
income.set_format('total income ($)', NumberFormatter(0)).show(5)
```


| ZIP | num returns | total income (\$) | num farmers |
|-------|-------------|-------------------|-------------|
| 90001 | 20970 | 531,772 | 0 |
| 90002 | 18960 | 467,128 | 0 |
| 90003 | 26180 | 618,848 | 0 |
| 90004 | 27360 | 1,649,431 | 0 |
| 90005 | 15430 | 717,290 | 0 |

... (1478 rows omitted)

Question 2.4. All ZIP codes with less than 100 returns (or some other special conditions) are grouped together into one ZIP code with a special code. Remove the row for that ZIP code from the `income` table.

Hint 1: This ZIP code value has far more returns than any of the other ZIP codes. Try using `group` and `sort` to find it.

Hint 2: To **remove** a row in the `income` table using `where`, assign `income` to the smaller table using the following expression structure:

```
income = income.where(...)
```

Hint 3: Each ZIP code is represented as a string, not an `int`.

```
In [10]: removeZIP = income.sort("num returns", descending=True)["ZIP"].item(0)
income = income.where("ZIP", are.not_containing(removeZIP))
income
```

```
Out[10]:
```

| ZIP | num returns | total income (\$) | num farmers |
|-------|-------------|-------------------|-------------|
| 90001 | 20970 | 531,772 | 0 |
| 90002 | 18960 | 467,128 | 0 |
| 90003 | 26180 | 618,848 | 0 |
| 90004 | 27360 | 1,649,431 | 0 |
| 90005 | 15430 | 717,290 | 0 |
| 90006 | 22630 | 571,157 | 0 |
| 90007 | 11710 | 315,581 | 0 |
| 90008 | 14710 | 668,523 | 0 |
| 90010 | 2210 | 320,471 | 0 |
| 90011 | 36670 | 864,961 | 0 |

... (1472 rows omitted)

Because each ZIP code has a different number of people, computing the average income across several ZIP codes requires some care. This will come up several times in this project. Here is a simple example:

Question 2.5 Among all the tax returns that

1. include a total income amount, and
2. are filed by people living in either ZIP code 94576 (a rural area north of Napa) or in ZIP code 94704 (a moderately-dense area in South Berkeley),

what is the average total income? **Express the answer in *dollars* as an `int` rounded to the nearest dollar.**

```
In [11]: # total income of the poeple living in 94576
total_income94576 = income.where("ZIP", are.equal_to("94576"))["total income"]
# tax return of the poeple living in 94576
tax_return94576 = income.where("ZIP", are.equal_to("94576"))["num returns"]
# total income of the people living in 94704
total_income94704 = income.where("ZIP", are.equal_to("94704"))["total income"]
# tax return of the poeple living in 94704
tax_return94704 = income.where("ZIP", are.equal_to("94704"))["num returns"]

average_income = round(sum(total_income94576+total_income94704)*1000/sum(tax_return94576+tax_return94704))
average_income
```

Out[11]: 52773

Question 2.6. Among all California tax returns that include a total income amount, what is the average total income? **Express the answer in *dollars* as an `int` rounded to the nearest dollar.**

```
In [12]: avg_total = round(sum(income["total income ($)"])*1000/sum(income["num returns"]))
avg_total
```

Out[12]: 72791

Farming


Farms use water, so it's plausible that farming is an important factor in water usage. Here, we will check for a relationship between farming and income.

Among the tax returns in California for ZIP codes represented in the `income` table, is there an association between income and living in a ZIP code with a lot of farmers?

We'll try to answer the question in 3 ways.

Question 2.7. Make a scatter plot with one point for each ZIP code. Display the average income *in dollars* on the vertical axis and the proportion of returns that are from farmers on the horizontal axis.

```
In [13]: # make a table containing the average income in dollars for each Zip code
avg_income = income.with_columns("average income", income["total income ($)"] / income["num returns"])
avg_income.scatter("proportion of farmers", "average income")
```

Question 2.8. From the graph, can you say whether ZIP codes with more farmers typically have lower or higher average income than ZIP codes with few or no farmers? Can you say how much lower or higher?

The range in the average income is greater in the areas with less proportion of farmers than the areas with greater proportion of farmers. The areas with the greater proportion of farmers have a steady average incomes. We cannot say that the ZIP codes with more farmers typically have lower or higher average income because the range in the average income is great in the areas with low number of farmers.

Question 2.9. Compare the average incomes for two groups of *tax returns*: those in ZIP codes with a greater-than-average proportion of farmers and those in ZIP codes with a less-than-average (or average) proportion. Make sure both of these values are displayed (preferably in a table). *Then, describe your findings.*

Hint: Make sure your result correctly accounts for the number of tax returns in each ZIP code, as in questions 2.5 and 2.6.

```
In [14]: # find the average proportion of farmers
farmer_avg_total = sum(income["num farmers"])/sum(income["num returns"])

# add a column of the average proportion of farmers in the table
income_farmers = income.with_column("average farmers", income["num farmers"]

# table with a greater-than-average proportion of farmers
farmers_greater = income_farmers.where("average farmers", are.above(farmer_avg_total))
# get the average incomes of tax returns in dollars rounded to the nearest c
```

```

income_avg_greater = round(sum(farmers_greater["total income ($)"])*1000/sum(
# table with a less-than-average proportion of farmers
farmers_less = income_farmers.where("average farmers", are.below_or_equal_to
# get the average incomes of tax returns in dollars rounded to the nearest c
income_avg_less = round(sum(farmers_less["total income ($)"])*1000/sum(farme

# Build and display a table with two rows:
# 1) incomes of returns in ZIP codes with a greater-than-average proportio
# 2) incomes of returns in other ZIP codes
avg_income_farmers = Table().with_columns("avg farmer proportion", ["above-a
"average income", [income_avg_great
avg_income_farmers


```

Out[14]: **avg farmer proportion average income**

| | |
|---------------|-------|
| above-average | 78231 |
| below-average | 71464 |

The areas with the average number of farmers greater than the average proportion have the higher average income than the areas with the average number of farmers below the average proportion.

Question 2.10. The graph below displays two histograms: the distribution of average incomes of ZIP codes that have above-average proportions of farmers, and that of ZIP codes with below-average proportions of farmers.



Are ZIP codes with below-average proportions of farmers more or less likely to have very low incomes? Explain how your answer is consistent with your answer to question 2.8.

The ZIP codes with below-average proportions of farmers more likely to have very low incomes than the ZIP codes with above-average proportions of farmers. However, the ZIP codes with below-average proportions of farmers have more likely to have high average incomes too. We can also see this from the graph in question 2.7. In the graph, the range of the average incomes are very large when the proportion of farmers are small, which is consistent with the graph in question 2.10.

ZIP codes cover all the land in California and do not overlap. Here's a map of all of them.



Question 2.11. Among the ZIP codes represented in the `income` table, is there an association between high average income and some aspect of the ZIP code's location? If so, describe one aspect of the location that is clearly associated with high income.

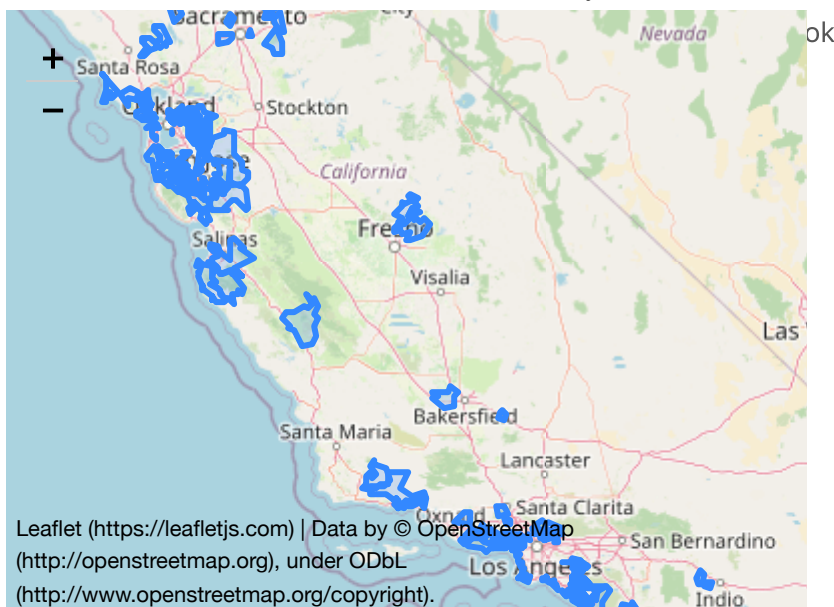
Answer the question by drawing a map of all ZIP codes that have an average income above 100,000 dollars. *Then, describe an association that you observe.*

In order to create a map of certain ZIP codes, you need to:

- Construct a table containing only the ZIP codes of interest, called `high_average_zips`.
- Join `high_average_zips` with the `zip_features` table to find the region for each ZIP code of interest.
- Call `Map(...)` on the column of features (provided).

```
In [15]: # Write code to draw a map of only the high-income ZIP codes.
# We have filled in some of it and suggested names for variables
# you might want to define.
zip_features = Table.from_records(zips.features)
high_average_zips = avg_income.where("average income", are.above(100000))
high_zips_with_region = zip_features.join('ZIP', high_average_zips)
Map(high_zips_with_region.column('feature'), width=400, height=300)
```

Out[15]:



Most of the ZIP codes with high income are gathered together near Los Angeles, San Jose, and San Francisco. These cities are located near coast and known to others.

Part 3: Water Usage

We will now investigate water usage in California. The `usage` table contains three columns:

- `PWSID` : The Public Water Supply Identifier of the district
- `Population` : Estimate of average population served in 2015
- `Water` : Average residential water use (gallons per person per day) in 2014-2015

```
In [16]: # Run this cell to create the usage table.

usage_raw.set_format(4, NumberFormatter)
max_pop = usage_raw.select(0, 1, 'population').group(0, max).relabel(2, 'P')
avg_water = usage_raw.select(0, 'res_gpcd').group(0, np.mean).relabel(1, 'W')
usage = max_pop.join('pwsid', avg_water).relabel(0, 'PWSID').relabel(1, 'Water')
usage
```

Out [16]:

| PWSID | District | Population | Water |
|---------|---|------------|---------|
| 0110001 | Alameda County Water District | 340000 | 70.7 |
| 0110003 | California Water Service Company Livermore | 57450 | 90.2727 |
| 0110005 | East Bay Municipal Utilities District | 1390000 | 76 |
| 0110006 | Hayward, City of | 151037 | 57.1818 |
| 0110008 | Pleasanton, City of | 73067 | 96.6364 |
| 0110009 | Dublin San Ramon Services District | 79547 | 68.6364 |
| 0110011 | Livermore, City of | 31994 | 85.8182 |
| 0310003 | Amador Water Agency | 23347 | 82.8182 |
| 0410002 | California Water Service Company Chico District | 101447 | 142 |
| 0410005 | California Water Service Company Oroville | 11208 | 88.8182 |

... (401 rows omitted)

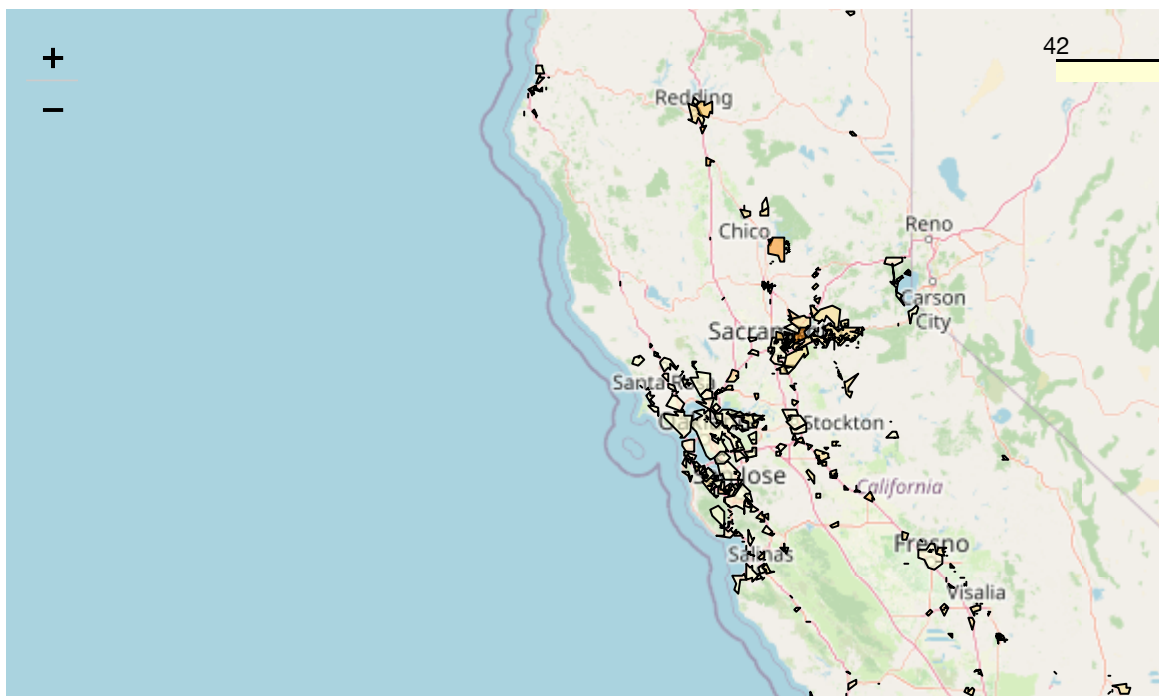
Question 3.1. Draw a map of the water districts, colored by the per capita water usage in each district.

Use the `districts.color(...)` method to generate the map. It takes as its first argument a two-column table with one row per district that has the district `PWSID` as its first column. The label of the second column is used in the legend of the map, and the values are used to color each region.

In [17]:

```
# We have filled in the call to districts.color(...). Set per_capita_usage
# to an appropriate table so that a map of all the water districts is
# displayed.
per_capita_usage = usage.select("PWSID","Water")
districts.color(per_capita_usage, key_on='feature.properties.PWSID')
```

Out [17]:



Question 3.2. Based on the map above, which part of California appears to use more water per person: the San Francisco area or the Los Angeles area?

The Los Angeles area use more water per person. There are more areas with the darker color in the Los Angeles area compared to the San Francisco area.

Next, we will try to match each ZIP code with a water district. ZIP code boundaries do not always line up with water districts, and one water district often covers multiple ZIP codes, so this process is imprecise. It is even the case that some water districts overlap each other. Nonetheless, we can continue our analysis by matching each ZIP code to the water district with the largest geographic overlap.

The table `wd_vs_zip` describes the proportion of land in each ZIP code that is contained in each water district and vice versa. (The proportions are approximate because they do not correctly account for discontinuous districts, but they're mostly accurate.)

In [18]: `wd_vs_zip.show(5)`

| PWSID | ZIP | District in ZIP | ZIP in District |
|---------|-------|-----------------|-----------------|
| 0110001 | 94536 | 9.41% | 68.51% |
| 0110001 | 94538 | 18.87% | 67.31% |
| 0110001 | 94539 | 13.13% | 44.36% |
| 0110005 | 94541 | 1.61% | 68.11% |
| 0110006 | 94541 | 18.68% | 98.46% |

... (3201 rows omitted)

Question 3.3. Complete the `district_for_zip` function that takes a ZIP code as its argument. It returns the PWSID with the largest value of `ZIP in District` for that `zip_code`, if that value is at least 50%. Otherwise, it returns the string `'No District'`.

```
In [19]: def district_for_zip(zip_code):
    zip_code = str(zip_code) # Ensure that the ZIP code is a string, not an
    districts = wd_vs_zip.where("ZIP", are.equal_to(zip_code)).sort("ZIP in
    at_least_half = districts.column("ZIP in District").item(0)
    if at_least_half >= 0.5:
        return districts["PWSID"].item(0)
    else:
        return 'No District'

district_for_zip(94709)
```

Out[19]: `'0110005'`

This function can be used to associate each ZIP code in the `income` table with a `PWSID` and discard ZIP codes that do not lie (mostly) in a water district.

```
In [20]: zip_pwsids = income.apply(district_for_zip, 'ZIP')
income_with_pwsid = income.with_column('PWSID', zip_pwsids).where('PWSID', is
income_with_pwsid.set_format(2, NumberFormatter(0)).show(5)
```

| ZIP | num returns | total income (\$) | num farmers | PWSID |
|-------|-------------|-------------------|-------------|---------|
| 90001 | 20970 | 531,772 | 0 | 1910067 |
| 90022 | 26680 | 767,484 | 0 | 1910036 |
| 90024 | 14690 | 4,395,487 | 20 | 1910067 |
| 90025 | 25110 | 4,019,082 | 20 | 1910067 |
| 90034 | 29950 | 1,828,572 | 0 | 1910067 |

... (662 rows omitted)

Question 3.4. Create a table called `district_data` with one row per PWSID and the following columns:

- `PWSID` : The ID of the district
- `Population` : Population estimate
- `Water` : Average residential water use (gallons per person per day) in 2014-2015
- `Income` : Average income in dollars of all tax returns in ZIP codes that are (mostly) contained in the district according to `income_with_pwsid`.

Hint: First create a `district_income` table that sums the incomes and returns for ZIP codes in each water district.

```
In [21]: district_income = income_with_pwsid.group("PWSID", np.sum).drop(1,4)

## add a column with the average income in dollars of all tax returns
district_income = district_income.with_column(
    "average income", district_income["total income ($) sum"]*1000/district_
district_income = district_income.drop(1,2)

district_data = usage.join("PWSID", district_income).relabelled(4, "Income")
district_data
```

```
Out[21]:
```

| PWSID | District | Population | Water | Income |
|---------|---|------------|---------|---------|
| 0110001 | Alameda County Water District | 340000 | 70.7 | 79032 |
| 0110005 | East Bay Municipal Utilities District | 1390000 | 76 | 82497.2 |
| 0110006 | Hayward, City of | 151037 | 57.1818 | 52923.5 |
| 0110008 | Pleasanton, City of | 73067 | 96.6364 | 163257 |
| 0110009 | Dublin San Ramon Services District | 79547 | 68.6364 | 133902 |
| 0410002 | California Water Service Company Chico District | 101447 | 142 | 50400.8 |
| 0410006 | South Feather Water and Power Agency | 18300 | 286.2 | 38720.5 |
| 0410011 | Del Oro Water Company | 9615 | 92.1818 | 44706.7 |
| 0710001 | Antioch, City of | 106455 | 110.273 | 53550.8 |
| 0710003 | Contra Costa Water District | 197536 | 101.636 | 73913.7 |

... (200 rows omitted)

Question 3.5. The `bay_districts` table gives the names of all water districts in the San Francisco Bay Area. Is there an association between water usage and income among Bay Area water districts? Use the tables you have created to compare water

usage between the 10 Bay Area water districts with the highest average income and the rest of the Bay Area districts, then describe the association. *Do not include any districts in your analysis for which you do not have income information.*

The names below are just suggestions; you may perform the analysis in any way you wish.

Note: Some Bay Area water districts may not appear in your `district_data` table. That's ok. Perform your analysis only on the subset of districts where you have both water usage & income information.

```
In [22]: bay_districts = Table.read_table('~DS_113_S23/Projects/Project_1/bay_district_data.csv')

## find the Bay Area districts
bay_district_data = district_data.where("District", are.contained_in(bay_districts["District"]))
## sort the district_data table for the top average income
bay_district_data = bay_district_data.sort("Income", descending=True)
bay_district_data.show(10)
```

| PWSID | District | Population | Water | Income |
|---------|---|------------|---------|-------------|
| 4110006 | California Water Service Company Bear Gulch | 58895 | 170.455 | 1.16049e+06 |
| 4310001 | California Water Service Company Los Altos/Suburban | 68163 | 124.545 | 349183 |
| 4310009 | Palo Alto, City of | 64403 | 89.1818 | 334057 |
| 4110017 | Menlo Park, City of | 16066 | 75.1818 | 278733 |
| 2110002 | Marin Municipal Water District | 188200 | 85.6364 | 176643 |
| 0110008 | Pleasanton, City of | 73067 | 96.6364 | 163257 |
| 4310007 | Mountain View, City of | 76781 | 65.5455 | 138570 |
| 4110021 | Estero Municipal Improvement District | 37165 | 66.3636 | 137893 |
| 4110001 | Mid-Peninsula Water District | 26730 | 83.7273 | 137537 |
| 4110008 | California Water Service Company Mid Peninsula | 135918 | 70.3636 | 135967 |

... (19 rows omitted)

Complete this one-sentence conclusion: In the Bay Area, people in the top 10 highest-income water districts used an average of **24.645** more gallons of water per person per day than people in the rest of the districts.

```
In [30]: # the average of water usage in the top 10 highest-income water districts
top10_avg_water = (sum(bay_district_data["Water"][0:10]))/10
# the average of water usage in the rest of the water districts
rest_avg_water = sum(bay_district_data["Water"][10:len(bay_district_data["Water"])]/len(bay_district_data["Water"])-10)

difference = top10_avg_water-rest_avg_water
difference
```

Out [30]: 24.645454545454541

Question 3.6. In one paragraph, summarize what you have discovered through the analyses in this project and suggest what analysis should be conducted next to better

understand California water usage, income, and geography. What additional data would be helpful in performing this next analysis?

When we see the maps, both the Los Angeles areas and San Francisco areas had the higher average income compared to other areas, but the Los Angeles areas have a higher water usage than the San Francisco areas. However, it is significant that the average income has a positive correlation with the water usage. For the better analysis, it would be good to note the number of farmers for each area, since the farmers use more water than the others and the areas with the higher average income usually did not had a lot of farmers.

Congratulations - you've finished Project 1 of Data 8!

To submit:

1. Select **Run All** from the **Cell** menu to ensure that you have executed all cells, including the test cells. Make sure that the visualizations you create are actually displayed.
2. Select **Download as HTML (.html)** from the **File** menu. (Sometimes that seems to fail. If it does, you can download as HTML, open the .html file in your browser, and print it to a PDF.)
3. Upload your HTML to Moodle,

If you want, draw some more maps below.

In [24]: *# Your extensions here (completely optional)*